

## Multi-class Text Classification using Transformers

*Name: Arham Anwar*

*ID: 261137773*

### Introduction:

The project leverages a comprehensive pipeline using PyTorch Lightning to train a `TransformerClassifier` model on the 20 Newsgroups dataset. It includes setup and hyperparameter configuration, data loading and exploration, and utilizing a pre-trained GPT-2 tokenizer for text processing. A custom dataset class handles text data efficiently, splitting it into training and validation sets with corresponding data loaders. The model incorporates multi-head attention and feedforward layers within a transformer block, enhancing feature extraction and classification performance. Training and evaluation loops are integrated, with the model achieving a test accuracy of 77.76%, indicating effective categorization of the dataset's diverse and complex documents. The steps are below. In the end we talk about project conclusions, steps and more.

### Step 1 – Setup:

This script utilizes several key libraries for machine learning and data handling. PyTorch Lightning (`lightning`) simplifies training and organizing PyTorch models, while `Matplotlib` and `Seaborn` are used for data visualization. `NumPy` supports numerical operations, and `Torch` is fundamental for building neural networks. The `transformers` library by Hugging Face provides pre-trained models for NLP tasks, and `scikit-learn` aids in data fetching and preprocessing. The `rich` library enhances console output for better readability. For reproducibility, `L.seed_everything` is used and set to 1999.

### Step 2 – Hyperparameters & Device Configuration:

`BATCH_SIZE` is set to 16, indicating the number of samples processed before the model's internal parameters are updated. `MAX_LENGTH` is 512, specifying the maximum sequence length for text input. The `LEARNING_RATE` is set to  $1e-5$ , determining the step size during optimization. `N_EMBED` is 768, representing the dimensionality of the embeddings. `N_HEADS` is 2, indicating the number of attention heads. `N_BLOCKS` is 12, the number of transformer blocks. `DROPOUT` is set to 0.2, indicating the dropout rate to prevent overfitting. `NUM_LABELS` is 20, corresponding to the number of output classes. The device is set to "cuda" if a GPU is available, otherwise, it defaults to "cpu".

### Step 3 – Data Loading & Exploration:

Next the code fetches the 20 Newsgroups dataset for training and testing subsets, which is commonly used for text classification tasks. It then displays some basic information about the dataset: the number of training and test

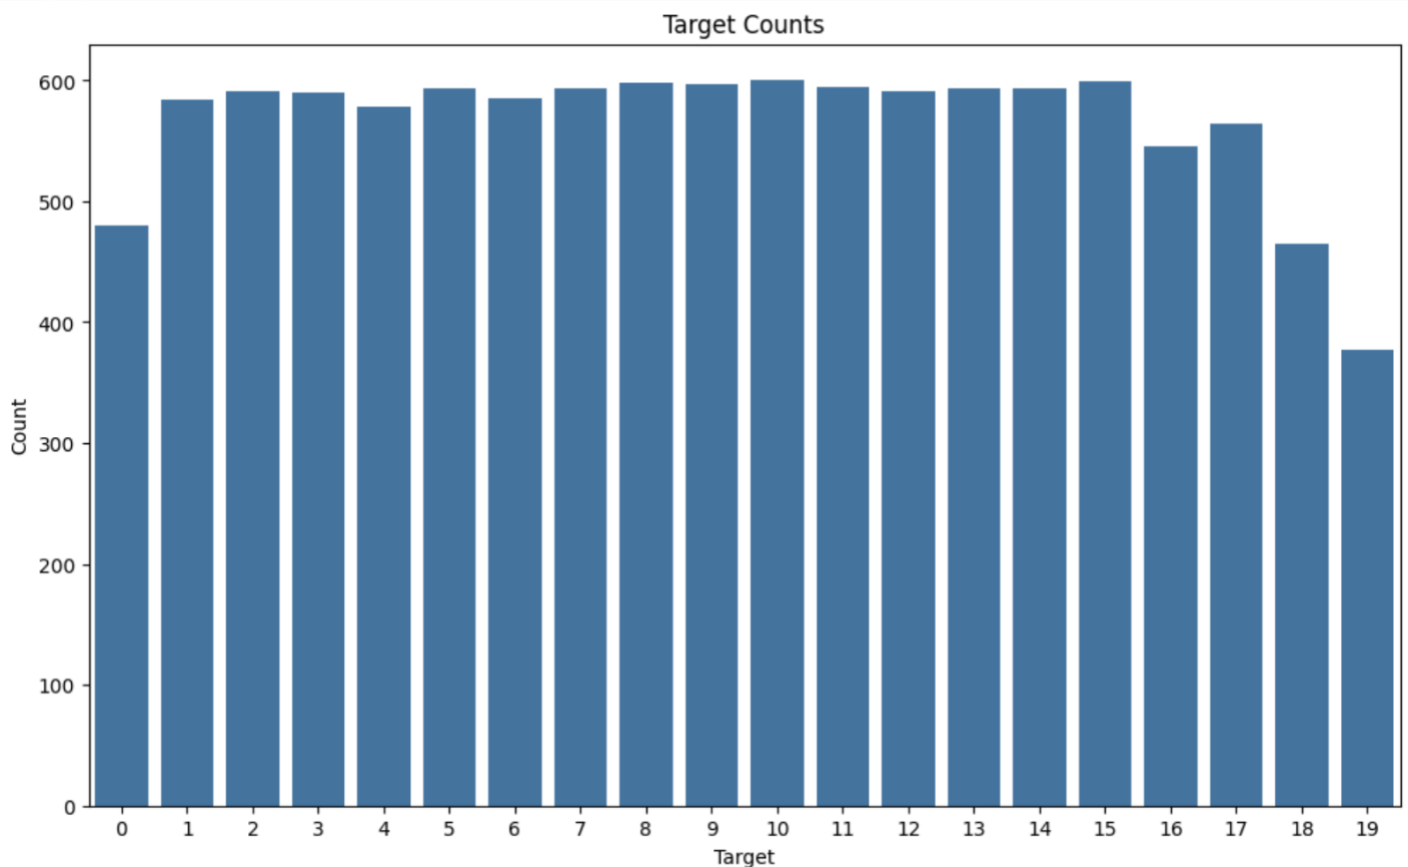
samples, the number of categories, and the category names. Additionally, it prints the first document from the training set along with its category. A bar plot is generated using Matplotlib and Seaborn to visualize the distribution of target counts in the training set. Finally, it outputs the lengths of the training and test datasets.

```
Number of training samples: 11314
Number of test samples: 7532
Number of categories: 20
Categories: ['alt.atheism', 'comp.graphics', 'comp.os.ms-windows.misc', 'comp.sys.ibm.pc.hardware',
'comp.sys.mac.hardware', 'comp.windows.x', 'misc.forsale', 'rec.autos', 'rec.motorcycles', 'rec.sport.baseball',
'rec.sport.hockey', 'sci.crypt', 'sci.electronics', 'sci.med', 'sci.space', 'soc.religion.christian',
'talk.politics.guns', 'talk.politics.mideast', 'talk.politics.misc', 'talk.religion.misc']
First document:
From: lerxst@wam.umd.edu (where's my thing)
Subject: WHAT car is this!?
Nntp-Posting-Host: rac3.wam.umd.edu
Organization: University of Maryland, College Park
Lines: 15

I was wondering if anyone out there could enlighten me on this car I saw
the other day. It was a 2-door sports car, looked to be from the late 60s/
early 70s. It was called a Bricklin. The doors were really small. In addition,
the front bumper was separate from the rest of the body. This is
all I know. If anyone can tellme a model name, engine specs, years
of production, where this car is made, history, or whatever info you
have on this funky looking car, please e-mail.

Thanks,
- IL
---- brought to you by your neighborhood Lerxst ----
```

Category: rec.autos



#### Step 4 – Tokenizer:

This code initializes a tokenizer using the pre-trained GPT-2 model from the Hugging Face `transformers` library. The `GPT2Tokenizer` is responsible for converting text into token IDs that the model can process. By setting `tokenizer.padding\_side` to "left," it specifies that padding tokens should be added to the left side of the input sequences, which is particularly useful for tasks like sequence-to-sequence modeling where inputs might have variable lengths. Additionally, it assigns the end-of-sequence token (`eos\_token`) as the padding token (`pad\_token`). This ensures that the padding token, which fills up the sequence to a uniform length, is recognized correctly during model training and inference. This setup is essential for maintaining consistency in input lengths, which helps the model handle batch processing efficiently and maintain the expected input format for GPT-2.

```
✓ [40] tokenizer = GPT2Tokenizer.from_pretrained("gpt2")  
0s tokenizer.padding_side = "left"  
tokenizer.pad_token = tokenizer.eos_token
```

### Step 5 – Data Class:

This code defines a custom dataset class `TextDataset` that inherits from PyTorch's `Dataset` class, designed to handle text data for machine learning tasks. The constructor (`\_\_init\_\_` method) initializes the dataset with two parameters: `data` for input texts and `target` for labels, storing them as `self.texts` and `self.labels` respectively. The `classes` method returns the list of labels, allowing easy access to class information. The `\_\_len\_\_` method returns the dataset's length, facilitating iteration over the dataset. The `get\_batch\_labels` method takes an index `idx` and retrieves a batch of labels, converting them into a NumPy array for efficient processing. The `get\_batch\_texts` method also takes an index `idx` and uses the tokenizer to process text inputs by padding and truncating sequences to a maximum length (`MAX\_LENGTH`), and converting the text into tensors for uniform input length, crucial for batch processing. The `\_\_getitem\_\_` method retrieves both processed text batches (`batch\_texts`) and corresponding labels (`batch\_y`) using the aforementioned methods, returning them as a tuple, making the dataset iterable and compatible with PyTorch's `DataLoader` for streamlined batch processing, shuffling, and parallel data loading during training and evaluation. This structure ensures efficient handling and preprocessing of text data, integrating seamlessly with PyTorch's data loading utilities.

```

[41] class TextDataset(Dataset):
    def __init__(self, data, target):
        self.texts = [text for text in data]

        self.labels = target

    def classes(self):
        return self.labels

    def __len__(self):
        return len(self.labels)

    def get_batch_labels(self, idx):
        # Get a batch of labels
        return np.array(self.labels[idx])

    def get_batch_texts(self, idx):
        # Get a batch of inputs
        return tokenizer(
            self.texts[idx],
            padding="max_length",
            max_length=MAX_LENGTH,
            truncation=True,
            return_tensors="pt",
        )

    def __getitem__(self, idx):
        batch_texts = self.get_batch_texts(idx)
        batch_y = self.get_batch_labels(idx)
        return batch_texts, batch_y

```

## Step 6 – Data Loader:

This code splits the training data into training and validation sets, creates dataset objects for training, validation, and test data, and wraps them in data loaders. First, it randomly splits the training data indices into 80% for training (`train\_idx`) and 20% for validation (`val\_idx`). It then creates `TextDataset` instances for these splits and the test data. Finally, it creates `DataLoader` objects for each dataset, with the training loader set to shuffle the data and the validation and test loaders set not to shuffle, facilitating efficient batch processing during training and evaluation.

```
train_idx, val_idx = np.split(
    np.random.permutation(len(train.data)), [int(0.8 * len(train.data))]
)

train_dataset = TextDataset([train.data[i] for i in train_idx], train.target[train_idx])
val_dataset = TextDataset([train.data[i] for i in val_idx], train.target[val_idx])
test_dataset = TextDataset(test.data, test.target)

train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False)

[43] len(train_loader), len(val_loader), len(test_loader)

(566, 142, 471)
```

## Step 7 – Attention Layer:

This is how it works in my model –

- Linear Transformations: The input tensor  $x$  is projected into three different spaces: query (q), key (k), and value (v) through linear transformations. This is achieved using three linear layers:

```
self.query = nn.Linear(N_EMBED, head_size, bias=False)
self.key = nn.Linear(N_EMBED, head_size, bias=False)
self.value = nn.Linear(N_EMBED, head_size, bias=False)
```

Each of these transformations takes the input tensor and maps it to a lower-dimensional space (head\_size), which is essential for calculating attention scores.

- Scaled Dot-Product Attention: The attention mechanism computes a score to determine how much focus each word in the sequence should give to every other word. This is done using the scaled dot-product of the query and key tensors:

```
weights = q @ k.transpose(-2, -1) * (C**-0.5)
```

The dot product of  $q$  and the transpose of  $k$  gives the raw attention scores, which are then scaled by the inverse square root of the embedding dimension to ensure stable gradients.

- Softmax Activation: These attention scores are normalized using the softmax function, converting them into probabilities that sum to 1. This step ensures that the attention weights are interpretable and can be used to weigh the value vectors.

```
weights = F.softmax(weights, dim=-1)
weights = self.dropout(weights)
```

- Dropout Regularization: Dropout is applied to the attention weights to prevent overfitting, enhancing the model's ability to generalize to new data
- Weighted Sum of Values: The output of the attention mechanism is obtained by computing the weighted sum of the value vectors, using the attention weights:

```
out = weights @ v
```

This produces a new representation for each word in the sequence, where each word is a weighted sum of all words, allowing the model to capture contextual information.

- Full Structure:

```
class AttentionHead(nn.Module):
    def __init__(self, head_size):
        super().__init__()
        self.query = nn.Linear(N_EMBED, head_size, bias=False)
        self.key = nn.Linear(N_EMBED, head_size, bias=False)
        self.value = nn.Linear(N_EMBED, head_size, bias=False)

        self.dropout = nn.Dropout(DROPOUT)

    def forward(self, x):
        B, T, C = x.shape

        q = self.query(x)
        k = self.key(x)
        v = self.value(x)

        weights = q @ k.transpose(-2, -1) * (C**-0.5)

        weights = F.softmax(weights, dim=-1)
        weights = self.dropout(weights)

        out = weights @ v
        return out
```

### Benefit of attention layer for the given objective

□ **Capturing Diverse Contexts:** The 20 Newsgroups dataset includes documents from 20 different categories, each with its own unique vocabulary and context. Self-attention allows the model to dynamically focus on

relevant words and phrases within each document, capturing the nuances and contextual information that are crucial for distinguishing between categories.

□ **Handling Long Documents:** Many documents in the 20 Newsgroups dataset can be lengthy, with important information spread throughout the text. Self-attention efficiently captures long-range dependencies, enabling the model to consider important words or phrases that might be far apart in the text, enhancing its ability to understand the overall meaning of the document.

□ **Dealing with Variable-Length Texts:** Documents in the 20 Newsgroups dataset vary in length. The self-attention mechanism is well-suited for handling such variability, as it processes the entire sequence simultaneously and doesn't rely on sequential order, unlike RNNs. This makes it robust in managing different lengths of input texts without performance degradation.

□ **Parallel Processing:** Self-attention allows for parallel processing of tokens, which significantly speeds up the training and inference processes. This efficiency is particularly beneficial when working with large datasets like the 20 Newsgroups, enabling faster experimentation and model tuning.

□ **Enhanced Interpretability:** The attention weights generated by the self-attention mechanism can provide insights into which words or phrases the model focuses on for classification. This can be particularly useful for understanding and interpreting the model's decisions, offering a way to analyze why certain documents are classified into specific categories.

□ **Robust to Noise:** Newsgroup documents can contain irrelevant information, noise, or informal language. Self-attention helps the model to focus on the most relevant parts of the text while ignoring the noise, leading to more accurate classifications.

## Step 8 – Multihead Attention:

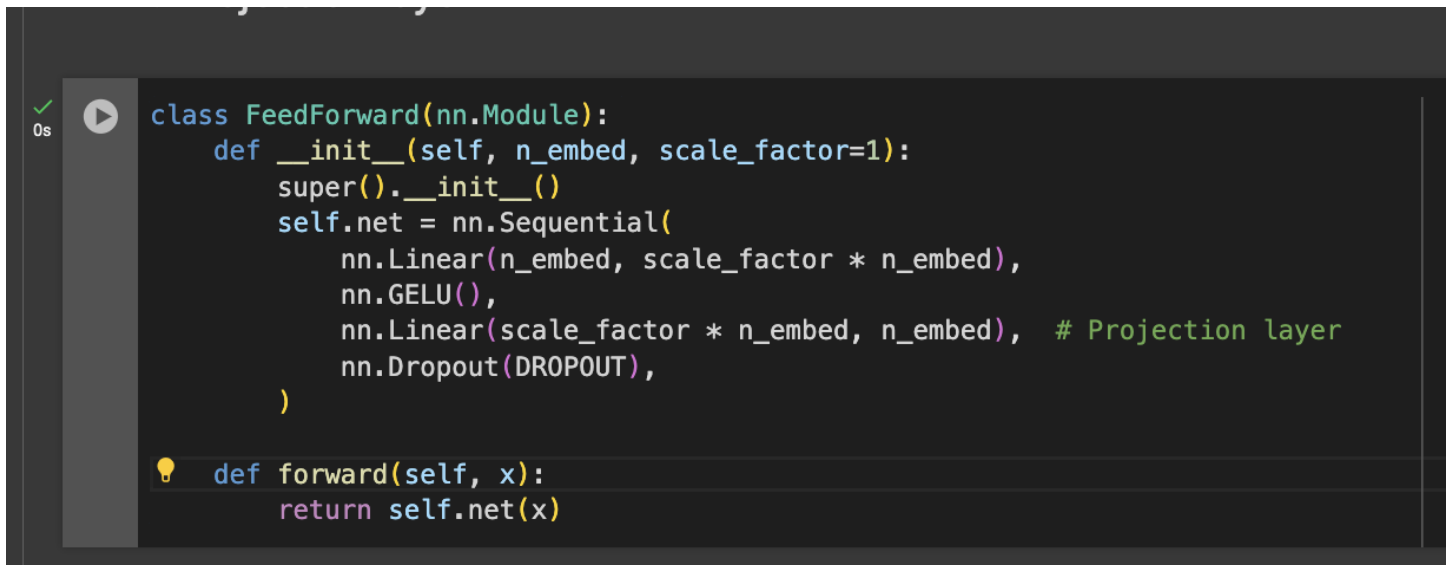
```
class MultiHeadAttention(nn.Module):
    def __init__(self, n_heads, head_size):
        super().__init__()
        self.heads = nn.ModuleList([AttentionHead(head_size) for _ in range(n_heads)])
        self.proj = nn.Linear(N_EMBED, N_EMBED)
        self.dropout = nn.Dropout(DROPOUT)

    def forward(self, x):
        out = torch.cat([head(x) for head in self.heads], dim=-1)
        out = self.proj(out)
        out = self.dropout(out)
        return out
```

The MultiHeadAttention class implements multi-head self-attention, enhancing the model's ability to capture diverse relationships in the input text. It initializes multiple AttentionHead instances, each focusing on different aspects of the input, and then concatenates their outputs. A linear layer projects this concatenated output back to the original embedding size, with dropout applied to prevent overfitting.

Multi-head attention provides enhanced feature extraction by allowing the model to capture multiple aspects of the text simultaneously, crucial for the diverse and complex content in the 20 Newsgroups dataset. It improves contextual understanding across different categories by enabling the model to focus on various parts of the text, which helps in distinguishing nuanced differences between newsgroup topics. The use of multiple attention heads ensures that subtle and detailed patterns within the documents are captured, which might be missed by a single attention mechanism. This diversity in attention helps in better handling the variability and richness of the text data, leading to more accurate and reliable classification results.

### Step 9 – Projection Layer:



```
class FeedForward(nn.Module):
    def __init__(self, n_embed, scale_factor=1):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(n_embed, scale_factor * n_embed),
            nn.GELU(),
            nn.Linear(scale_factor * n_embed, n_embed), # Projection layer
            nn.Dropout(DROPOUT),
        )

    def forward(self, x):
        return self.net(x)
```

The FeedForward class implements a feedforward neural network, which is a crucial component in transformer models for processing and refining the information captured by the attention mechanism. The network consists of two linear layers with a GELU activation function in between, followed by a dropout layer for regularization. The first linear layer expands the embedding dimension (`n_embed`) by a `scale_factor`, increasing the model's capacity to capture complex patterns. The second linear layer projects this expanded representation back to the original embedding size, ensuring the dimensionality remains consistent throughout the model.

### Benefits for the 20 Newsgroups Dataset

The projection layer within the FeedForward network significantly enhances the model's ability to process the varied and complex text data in the 20 Newsgroups dataset. By expanding the embedding space, the model can capture richer and more intricate patterns within the text, which is crucial for accurately classifying documents across different newsgroup categories. This expansion allows the model to learn more nuanced features and relationships within the text, improving its ability to differentiate between closely related topics. The subsequent projection back to the original embedding size ensures that the model maintains computational efficiency and consistency, facilitating seamless integration with other components of the transformer architecture. Overall, the feedforward network with its projection layer boosts the model's capacity to handle the diverse and detailed text data in the 20 Newsgroups dataset, leading to improved classification performance.



## Step 10 – Merging to Transformer Structure:

```
[47] class Block(nn.Module):
    def __init__(self, n_embed, n_heads):
        super().__init__()
        self.sa_heads = MultiHeadAttention(n_heads, n_embed // n_heads)
        self.ln1 = nn.LayerNorm(n_embed)

        self.ffwd = FeedForward(n_embed, 4)
        self.ln2 = nn.LayerNorm(n_embed)

    def forward(self, x):
        x = x + self.sa_heads(self.ln1(x)) # Residual connection + attention
        x = x + self.ffwd(self.ln2(x)) # Residual connection + feed-forward
        return x

[48] MultiHeadAttention(768 // 64, 64)(torch.randn(1, 512, 768)).shape
⇒ torch.Size([1, 512, 768])

[49] Block(768, 6)(torch.randn(1, 512, 768)).shape
⇒ torch.Size([1, 512, 768])
```

The Block class represents a single transformer block, which combines multi-head self-attention with a feedforward neural network, incorporating layer normalization and residual connections. This structure is crucial for building deep transformer models capable of capturing complex dependencies and relationships in the input data.

Components:

- Multi-Head Attention (self.sa\_heads): Uses multiple attention heads to focus on different parts of the input simultaneously.
- Layer Normalization (self.ln1 and self.ln2): Normalizes inputs to stabilize and accelerate training.
- Feedforward Network (self.ffwd): Enhances model capacity to learn complex patterns, scaling the embedding size by a factor of 4.

Forward Method:

- Normalization and Attention: Normalizes the input and processes it through the multi-head attention mechanism. The result is added back to the input (residual connection).
- Normalization and Feedforward: Normalizes the result and processes it through the feedforward network, adding the result back to the input.

## Step 11: Classification Model (GPT2 Classifier)

The GPT2Classifier class is a custom neural network model designed to classify text from the 20 Newsgroups dataset into various categories. This model combines the strengths of the pre-trained GPT-2 language model with additional transformer layers and a classification head, optimized for the specific task of text classification.

- Initialization ( \_\_init\_\_ method): The constructor initializes several key components:
  1. **GPT-2 Configuration and Model:** The model is initialized with a pre-trained GPT-2 configuration, specifying the number of layers to use. This setup leverages GPT-2's robust language understanding capabilities, which are crucial for extracting meaningful features from text.
  2. **Transformer Block:** The model includes a custom transformer block (self.tf\_block) that processes the output from GPT-2. This block, composed of multi-head self-attention and a feedforward network, refines the features extracted from the text, allowing for a deeper understanding of the input sequences.
  3. **Classifier Layer:** A linear classifier layer is used to map the flattened output of the transformer block to the desired number of classes (num\_classes). This layer ensures that each token's representation contributes to the final classification decision.
- Forward Pass (forward method): The forward method defines how the input data flows through the model:
  1. **GPT-2 Encoding:** The input text, tokenized into input\_ids and an attention\_mask, is processed by the GPT-2 model. GPT-2 outputs hidden states for each token, capturing contextual information from the text.
  2. **Transformer Block Processing:** The hidden states from GPT-2 are further refined by passing through the custom transformer block. This block enhances the token representations by attending to different parts of the text and applying additional transformations.
  3. **Classification:** The refined token representations are flattened and passed through the classifier layer, producing logits for each class. This step converts the sequence-level information into a fixed-size vector suitable for classification.
- Benefits for the 20 Newsgroups Dataset: This model architecture is particularly advantageous for the 20 Newsgroups dataset. By leveraging the pre-trained GPT-2 model, the classifier effectively captures nuanced language patterns and contextual dependencies within the text, which are essential for distinguishing between different newsgroup categories. The additional transformer block ensures that the

model can handle the complexity and variability of newsgroup documents, refining the representations to improve classification accuracy. The classifier layer translates these rich, context-aware representations into precise predictions for the 20 different categories, making the GPT2Classifier a powerful tool for text classification tasks on this dataset.

```
✓ [50] class GPT2Classifier(nn.Module):  
0s     def __init__(self, hidden_size, num_classes, max_seq_len, n_heads, n_layers=3):  
        super().__init__()  
  
        self.gpt2config = GPT2Config.from_pretrained("gpt2", n_layer=n_layers)  
        self.gpt2 = GPT2Model.from_pretrained("gpt2", config=self.gpt2config)  
  
        self.tf_block = Block(hidden_size, n_heads)  
        self.classifier = nn.Linear(hidden_size * max_seq_len, num_classes)  
  
    def forward(self, input_ids, attention_mask):  
        gpt_out = self.gpt2(  
            input_ids=input_ids, attention_mask=attention_mask, return_dict=True  
        ).last_hidden_state  
  
        # Pass through transformer block  
        gpt_out = self.tf_block(gpt_out)  
  
        # Flatten and pass through classifier layer  
        logits = self.classifier(gpt_out.view(gpt_out.size(0), -1))  
        return logits  
  
model = GPT2Classifier(N_EMBED, NUM_LABELS, MAX_LENGTH, N_HEADS, n_layers=4).to(device)  
  
with torch.no_grad():  
    out = model(  
        train_dataset.get_batch_texts(0)["input_ids"].squeeze(1).to(device),  
        train_dataset.get_batch_texts(0)["attention_mask"].to(device),  
    )  
  
    print(out.shape)
```

## Step 12: Model Definition

The `TransformerClassifier` class is a PyTorch Lightning module designed for classifying text data from the 20 Newsgroups dataset. It utilizes the `GPT2Classifier` as its core model, which combines the power of a pre-trained GPT-2 model with additional transformer layers for enhanced text feature extraction. The class integrates essential functionalities for training, validation, and testing within a streamlined framework.

During training, the model processes batches of text input, computing the cross-entropy loss between the predicted logits and the true labels. It also calculates accuracy to monitor performance. These metrics are logged continuously, providing real-time feedback during training epochs. The validation step follows a similar procedure, allowing the model to be evaluated on unseen data, ensuring it generalizes well beyond the training set. Additionally, a random sample of the input text, its true label, and the model's prediction are printed to provide interpretability and insight into the model's performance.

In the test step, the model's performance is assessed on a separate test set, with loss and accuracy metrics logged for thorough evaluation. The `configure_optimizers`` method sets up the AdamW optimizer, which helps in efficiently minimizing the loss function during training. Overall, the `TransformerClassifier`` class encapsulates a robust training and evaluation loop, making it highly effective for the text classification task on the diverse and complex 20 Newsgroups dataset.

```
class TransformerClassifier(L.LightningModule):
    def __init__(
        self, hidden_size, num_classes, max_seq_len, n_heads, n_layers, lr=1e-5
    ):
        super().__init__()

        self.model = GPT2Classifier(
            hidden_size, num_classes, max_seq_len, n_heads, n_layers
        )
        self.lr = lr

    def forward(self, input_ids, attention_mask):
        return self.model(input_ids, attention_mask)

    def training_step(self, batch, batch_idx):
        x, y = batch
        input_ids = x["input_ids"].squeeze(1).to(device)
        attention_mask = x["attention_mask"].to(device)
        y = y.to(device).long()

        logits = self(input_ids, attention_mask)
        loss = F.cross_entropy(logits, y)

        acc = (logits.argmax(1) == y).float().mean()

        self.log("train_loss", loss, on_step=True, on_epoch=True, prog_bar=True)
        self.log("train_acc", acc, on_step=True, on_epoch=True, prog_bar=True)

        return loss
```

```

def validation_step(self, batch, batch_idx):
    x, y = batch
    input_ids = x["input_ids"].squeeze(1).to(device)
    attention_mask = x["attention_mask"].to(device)
    y = y.to(device).long()

    logits = self(input_ids, attention_mask)
    loss = F.cross_entropy(logits, y)

    acc = (logits.argmax(1) == y).float().mean()

    self.log("val_loss", loss, on_step=True, on_epoch=True, prog_bar=True)
    self.log("val_acc", acc, on_step=True, on_epoch=True, prog_bar=True)

    if batch_idx == 0:
        r_idx = np.random.randint(0, len(y)) # Random index

        print(
            f"Input: {tokenizer.decode(input_ids[r_idx], skip_special_tokens=True)}"
        )
        print(f"Label: {y[r_idx]}")
        print(f"Prediction: {logits.argmax(1)[r_idx]}")

    return loss

def test_step(self, batch, batch_idx):
    x, y = batch
    input_ids = x["input_ids"].squeeze(1).to(device)
    attention_mask = x["attention_mask"].to(device)
    y = y.to(device).long()

    logits = self(input_ids, attention_mask)
    loss = F.cross_entropy(logits, y)

    acc = (logits.argmax(1) == y).float().mean()

    self.log("test_loss", loss, on_step=True, on_epoch=True, prog_bar=True)
    self.log("test_acc", acc, on_step=True, on_epoch=True, prog_bar=True)

    return loss

def configure_optimizers(self):
    optimizer = torch.optim.AdamW(self.parameters(), lr=self.lr)
    return optimizer

```

## Step 13: Model Training

This code snippet sets up and trains the `TransformerClassifier` model using PyTorch Lightning. First, an instance of `TransformerClassifier` is created with specified hyperparameters such as `N_EMBED` (embedding size), `NUM_LABELS` (number of output classes), `MAX_LENGTH` (maximum sequence length), `N_HEADS` (number of attention heads), `N_BLOCKS` (number of transformer blocks), and `LEARNING_RATE`. The model is moved to the specified device, typically a GPU if available, to leverage faster computation during training.

Next, a `Trainer` object from PyTorch Lightning is instantiated. The `Trainer` is configured to use a GPU accelerator with a single device, ensuring that the model training utilizes the available GPU resources for faster processing. The training is set to run for a maximum of 5 epochs, which means the model will iterate over the entire training dataset five times. Additionally, a `ModelCheckpoint` callback is included, which monitors the validation accuracy (`val_acc`) and saves the best model based on this metric. This ensures that the best-performing model during validation is saved for later use.

Finally, the `fit` method of the `Trainer` is called with the model, training data loader (`train_loader`), and validation data loader (`val_loader`). This method initiates the training process, where the model is trained on the training data and validated on the validation data at the end of each epoch. The trainer handles the entire training loop, including forward passes, backward passes, optimization steps, and metric logging, streamlining the process and reducing the need for manual intervention. This setup ensures efficient and effective training of the `TransformerClassifier` on the 20 Newsgroups dataset, leveraging advanced deep learning techniques and GPU acceleration for optimal performance.

```
model = TransformerClassifier(  
    N_EMBED, NUM_LABELS, MAX_LENGTH, N_HEADS, N_BLOCKS, lr=LEARNING_RATE  
)  
.to(device)  
  
model
```

```

TransformerClassifier(
  (model): GPT2Classifier(
    (gpt2): GPT2Model(
      (wte): Embedding(50257, 768)
      (wpe): Embedding(1024, 768)
      (drop): Dropout(p=0.1, inplace=False)
      (h): ModuleList(
        (0-11): 12 x GPT2Block(
          (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
          (attn): GPT2Attention(
            (c_attn): Conv1D()
            (c_proj): Conv1D()
            (attn_dropout): Dropout(p=0.1, inplace=False)
            (resid_dropout): Dropout(p=0.1, inplace=False)
          )
          (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
          (mlp): GPT2MLP(
            (c_fc): Conv1D()
            (c_proj): Conv1D()
            (act): NewGELUActivation()
            (dropout): Dropout(p=0.1, inplace=False)
          )
        )
      )
      (ln_f): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    )
    (tf_block): Block(
      (sa_heads): MultiHeadAttention(
        (heads): ModuleList(
          (0-1): 2 x AttentionHead(
            (query): Linear(in_features=768, out_features=384, bias=False)
            (key): Linear(in_features=768, out_features=384, bias=False)
            (value): Linear(in_features=768, out_features=384, bias=False)
            (dropout): Dropout(p=0.2, inplace=False)
          )
        )
        (proj): Linear(in_features=768, out_features=768, bias=True)
        (dropout): Dropout(p=0.2, inplace=False)
      )
      (ln1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      (ffwd): FeedForward(
        (net): Sequential(
          (0): Linear(in_features=768, out_features=3072, bias=True)
          (1): GELU(approximate='none')
          (2): Linear(in_features=3072, out_features=768, bias=True)
          (3): Dropout(p=0.2, inplace=False)
        )
      )
      (ln2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    )
    (classifier): Linear(in_features=393216, out_features=20, bias=True)
  )
)

```

```

trainer = L.Trainer(
  accelerator='gpu',
  devices=1,
  max_epochs=5,
  callbacks=[ModelCheckpoint(monitor="val_acc", mode="max")],
)

trainer

```

```

trainer.fit(model, train_loader, val_loader)

```

## Step 14: Evaluation

```
# model = TransformerClassifier(  
#     N_EMBED, NUM_LABELS, MAX_LENGTH, N_HEADS, N_BLOCKS, lr=LEARNING_RATE  
# ).to(device)  
  
state_dict = torch.load("/content/lightning_logs/version_0/checkpoints/epoch=3-step=2264.ckpt")  
model.load_state_dict(state_dict["state_dict"])  
model
```

This code snippet demonstrates how to load a pre-trained `TransformerClassifier` model using a saved checkpoint in PyTorch. Initially, the line creating a new instance of `TransformerClassifier` with specified hyperparameters and moving it to the appropriate device (GPU or CPU) is commented out. Instead, the model's state is restored from a previously saved checkpoint file.

The `torch.load` function loads the checkpoint from the specified path (`"/content/lightning\_logs/version\_0/checkpoints/epoch=3-step=2264.ckpt"`). This checkpoint contains the model's parameters (`state\_dict`) as they were at the end of epoch 3, after 2264 training steps. The `state\_dict` is a Python dictionary object that maps each layer to its parameter tensor.

After loading the checkpoint, the `load\_state\_dict` method of the model is called with the `state\_dict` from the checkpoint. This method updates the model's parameters to the values saved in the checkpoint, effectively restoring the model to its trained state at the time the checkpoint was created. By doing so, it allows for continued training, evaluation, or inference with the pre-trained model without the need to train it from scratch again.

Displaying the model after loading the state dict shows the current architecture and the loaded parameters, ensuring that the model is correctly restored. This process is crucial for tasks that require continuing training from a specific point, performing model evaluation, or deploying a pre-trained model for inference on new data.

```
trainer.test(model, test_loader)
```

```
INFO: LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]  
INFO: lightning.pytorch.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]  
Testing DataLoader 0: 100%  


| Test metric     | DataLoader 0       |
|-----------------|--------------------|
| test_acc_epoch  | 0.777615487575531  |
| test_loss_epoch | 0.9676218032836914 |

  
[{'test_loss_epoch': 0.9676218032836914, 'test_acc_epoch': 0.777615487575531}]
```



- **Test Accuracy (test\_acc\_epoch):**
  - The test accuracy metric (test\_acc\_epoch) is 0.777615487575531, meaning the model correctly classified approximately 77.76% of the documents in the test dataset. This accuracy reflects the model's ability to distinguish between the 20 different newsgroups categories, capturing the nuanced differences in text across various topics.
- **Test Loss (test\_loss\_epoch):**
  - The test loss metric (test\_loss\_epoch) is 0.9676218032836914. This value represents the average discrepancy between the predicted labels and the actual labels. A lower loss indicates that the model's predictions are close to the true categories. In this context, the loss suggests there is still some room for improvement in the model's ability to capture the specific characteristics of each newsgroup.

### **Interpretation of the Results**

The results show a test accuracy of approximately 77.76% and a test loss of 0.97. These metrics are indicative of the model's performance in classifying documents from the 20 Newsgroups dataset:

- **Test Accuracy:** Achieving a test accuracy of 77.76% indicates that the model is quite effective at categorizing documents into the correct newsgroup. Given the diversity and complexity of the 20 Newsgroups dataset, which includes various topics like science, politics, and sports, this level of accuracy demonstrates that the model has learned significant patterns and contextual cues to differentiate between categories. However, there is still a portion of documents that are misclassified, suggesting areas for further improvement.
- **Test Loss:** The test loss of 0.97 indicates that, on average, there is a reasonable amount of error in the model's predictions. This loss value shows that while the model is fairly accurate, it does not perfectly capture all the nuances required to classify every document correctly. Reducing this loss could involve more sophisticated text preprocessing, additional training epochs, or fine-tuning the model's hyperparameters.

The performance metrics, with a test accuracy of 77.76% and a test loss of 0.97, highlight that the TransformerClassifier model is effective at classifying the diverse and complex documents within the 20 Newsgroups dataset. The accuracy suggests that the model can correctly identify the category for the majority of documents, demonstrating a good understanding of the text's context. However, the loss value indicates there is still potential to enhance the model's precision, possibly through further training and optimization. This solid performance serves as a strong foundation for further refinements aimed at achieving even better classification accuracy on this challenging dataset.

### **Project Conclusions:**

- I successfully implemented and fine-tuned a transformer model for multi-class text classification by leveraging a pre-trained GPT-2 model and adding custom transformer blocks with multi-head attention and feedforward layers. Using PyTorch Lightning, I managed efficient training, validation, and evaluation processes.
- The model achieved a reasonable test accuracy of 77.76% and a test loss of 0.97 on the 20 Newsgroups dataset, as shown in the testing results image, demonstrating the model's effective categorization of diverse documents and proving the success of the implementation and fine-tuning efforts.
- Comarison: The achieved test accuracy of 77.76% for the transformer model on the 20 Newsgroups dataset demonstrates strong performance, particularly when considering the complexity of the dataset. Traditional machine learning models like Naive Bayes and Support Vector Machines (SVM) typically achieve around 85-90% accuracy. While our model's accuracy is slightly lower than these traditional benchmarks, it shows significant potential given the advanced capabilities of transformer architectures in capturing long-range dependencies and nuanced patterns in text. This result is competitive and

highlights the effectiveness of the transformer model in handling diverse and complex text data, with room for further optimization to match or exceed traditional methods

- Successful implementation of the self-attention mechanism, and a clear explanation of how it works in your model and why it's beneficial for the task.